

Introduction

Numerical Methods for Deep Learning

Course Overview

- ▶ Module 1: Linear Models
 - 2. Linear Models and Least-Squares
 - 3. Iterative Methods for Least-Squares
 - 4. Linear Models for Classification
 - 5. Newton's Method for Classification
 - 6. Regularization for Image Classification

Course Overview

- ▶ Module 2: Neural Networks
 - 7. Introduction to Nonlinear Models
 - 8. Single Layer Neural Networks
 - 9. Training Algorithms for Single Layer Neural Networks
 - 10. Introduction to Deep Neural Networks
 - 11. Differentiating Deep Neural Networks
 - 12. Stochastic Gradient Descent and Variants
- ▶ Module 3: Parametric Models/Convolution Neural Networks
 - 13. Introduction to Parametric Models
 - 14. Application of CNN: Image Segmentation
 - 15. CNN and their relation to PDEs

Neural Networks: History

- ▶ Neural Networks with a particular (deep) architecture
- ▶ Exist for a long time (70's and even earlier) [10, 11, 8]
- ▶ Recent revolution - computational power and lots of data [1, 9, 7]
- ▶ Can perform very well for large amounts of data
- ▶ Applications
 - ▶ Image recognition [4, 6, 7], segmentation, natural language processing [2, 3, 5]

Neural Networks: History

- ▶ Neural Networks with a particular (deep) architecture
- ▶ Exist for a long time (70's and even earlier) [10, 11, 8]
- ▶ Recent revolution - computational power and lots of data [1, 9, 7]
- ▶ Can perform very well for large amounts of data
- ▶ Applications
 - ▶ Image recognition [4, 6, 7], segmentation, natural language processing [2, 3, 5]
- ▶ A few recent news articles:
 - ▶ Apple Is Bringing the AI Revolution to Your iPhone, WIRED 2016
 - ▶ Why Deep Learning Is Suddenly Changing Your Life, FORTUNE 2016
 - ▶ Data Scientist: Sexiest Job of the 21st Century, Harvard Business Rev 17

NN - A Quick Overview

Neural Networks is a data interpolator/classifier when the underlying model is unknown.

A generic way to write it is

$$\mathbf{c} = f(\mathbf{y}, \boldsymbol{\theta}).$$

- ▶ The function f is the computational model.
- ▶ $\mathbf{y} \in \mathbb{R}^{n_f}$ is the input data (e.g., an image)
- ▶ $\mathbf{c} \in \mathbb{R}^{n_c}$ is the output (e.g. class the image)
- ▶ $\boldsymbol{\theta} \in \mathbb{R}^{n_p}$ are parameters of the model f

In learning we have examples $\{(\mathbf{y}_j, \mathbf{c}_j) : j = 1, \dots, n\}$ and the goal is to estimate or “learn” the parameters $\boldsymbol{\theta}$

Learning From Data: The Core of Science

Given inputs and outputs, how to choose f ?

Option 1 (Fundamental(?) understanding): For example, Newton's formula

$$x(t) = \frac{1}{2}gt^2,$$

with unknown parameter g .

Learning From Data: The Core of Science

Given inputs and outputs, how to choose f ?

Option 1 (Fundamental(?) understanding): For example, Newton's formula

$$x(t) = \frac{1}{2}gt^2,$$

with unknown parameter g .

To estimate g observe falling object

t	x
0	0
1	4.9
2	20.1
3	44.1

What is the optimal value for g ?

Learning From Data: The Core of Science

Given inputs and outputs, how to choose f ?

Option 2 (Phenomenological models): For example, Archie's law - what is the electrical resistivity of a rock and how it relates to its porosity, ϕ and saturation, S_w ?

$$\rho(\phi, S_w) = a\phi^{n/2}S_w^p$$

a, n, p unknown parameters

Obtaining parameters from observed data and lab experiments on rocks

Phenomenological vs. Fundamental

Fundamental laws come from understanding(?) the underlying process. They are **assumed invariant** and can therefore be predictive(?).

Phenomenological models are data driven. They “work” on some given data. Hard to know what are the limitations.

But ...

- ▶ models based on understanding can do poorly - weather, economics ...
- ▶ models based on data can sometimes do better
- ▶ how do we quantify understanding?

Generalization

Suppose that we have examples $\{\mathbf{y}_j, \mathbf{c}_j\}$, $j = 1, \dots, n$, a model $f(\mathbf{y}, \boldsymbol{\theta})$ and some optimal parameter $\boldsymbol{\theta}^*$.

Let $\{(\mathbf{y}_j^t, \mathbf{c}_j^t) : j = 1, \dots, s\}$ be some test set, that was not used to compute $\boldsymbol{\theta}^*$.

Generalization

Suppose that we have examples $\{\mathbf{y}_j, \mathbf{c}_j\}$, $j = 1, \dots, n$, a model $f(\mathbf{y}, \boldsymbol{\theta})$ and some optimal parameter $\boldsymbol{\theta}^*$.

Let $\{(\mathbf{y}_j^t, \mathbf{c}_j^t) : j = 1, \dots, s\}$ be some test set, that was not used to compute $\boldsymbol{\theta}^*$.

Loosely speaking, if

$$\|f(\mathbf{y}_j^t, \boldsymbol{\theta}^*) - \mathbf{c}_j^t\|_p \text{ is small}$$

then the model is predictive - it generalizes well

Generalization

Suppose that we have examples $\{\mathbf{y}_j, \mathbf{c}_j\}$, $j = 1, \dots, n$, a model $f(\mathbf{y}, \boldsymbol{\theta})$ and some optimal parameter $\boldsymbol{\theta}^*$.

Let $\{(\mathbf{y}_j^t, \mathbf{c}_j^t) : j = 1, \dots, s\}$ be some test set, that was not used to compute $\boldsymbol{\theta}^*$.

Loosely speaking, if

$$\|f(\mathbf{y}_j^t, \boldsymbol{\theta}^*) - \mathbf{c}_j^t\|_p \text{ is small}$$

then the model is predictive - it generalizes well

For phenomenological models, there is no reason why the model should generalize, but in practice it often does.

Generalization

Why would a model generalize poorly?

$$1 \ll \|f(\mathbf{y}_j^t, \boldsymbol{\theta}^*) - \mathbf{c}_j^t\|_p$$

Generalization

Why would a model generalize poorly?

$$1 \ll \|f(\mathbf{y}_j^t, \boldsymbol{\theta}^*) - \mathbf{c}_j^t\|_p$$

Two common reasons:

1. Our “optimal” $\boldsymbol{\theta}^*$ was optimal for the training but is less so for other data
2. The chosen computational model f is poor (e.g. linear model for a nonlinear function).

Example: Classification of Hand-written Digits

- ▶ Let $\mathbf{y}_j \in \mathbb{R}^{n_f}$ and let $\mathbf{c}_j \in \mathbb{R}^{n_c}$.
- ▶ The vector \mathbf{c} is the probability of \mathbf{y} belonging to a certain class. Clearly, $0 \leq \mathbf{c}_j \leq 1$ and $\sum_{j=1}^{n_c} \mathbf{c}_j = 1$.

Examples (MNIST):

\mathbf{y}_1



\mathbf{y}_2



$$\mathbf{c}_1 = [0, 0, 0, 0, 1, 0, 0, 0, 0, 0]^\top \quad \mathbf{c}_2 = [0, 0.3, 0, 0, 0, 0, 0, 0, 0.7, 0]^\top$$

Example: Classification of Natural Images

Image classification of natural images

Examples (CIFAR-10):



Example: Semantic Segmentation

- ▶ let $\mathbf{y}_j \in \mathbb{R}^n$ be an RGB or grey valued image.
- ▶ let the pixels in $\mathbf{c}_j \in \{1, 2, 3, \dots\}^k$ denote the labels.

\mathbf{y} , input image



\mathbf{c} , segmentation (labeled image)



Goal: Find map $\mathbf{c} = f(\mathbf{y}, \theta)$

Example 3: Semantic Segmentation

Problem: Given image \mathbf{y} and label \mathbf{c} find a map $f(\cdot, \boldsymbol{\theta})$ such that $\mathbf{c} \approx f(\mathbf{y}, \boldsymbol{\theta})$

Example 3: Semantic Segmentation

Problem: Given image \mathbf{y} and label \mathbf{c} find a map $f(\cdot, \boldsymbol{\theta})$ such that $\mathbf{c} \approx f(\mathbf{y}, \boldsymbol{\theta})$

First step: Reduce the dimensionality of problem.

- ▶ extract features from the image
- ▶ classify in the feature space

Reduce the problem of learning from the image to feature detection and classification

Example 3: Semantic Segmentation

Problem: Given image \mathbf{y} and label \mathbf{c} find a map $f(\cdot, \boldsymbol{\theta})$ such that $\mathbf{c} \approx f(\mathbf{y}, \boldsymbol{\theta})$

First step: Reduce the dimensionality of problem.

- ▶ extract features from the image
- ▶ classify in the feature space

Reduce the problem of learning from the image to feature detection and classification

Possible features: Color, neighbors, edges ...


Example 3 - Semantic Segmentation

Simpler setup

- ▶ data, \mathbf{y} is the RGB value of the pixel (and its neighbors?)
- ▶ \mathbf{c} is a labeled pixel
- ▶ The map $\mathbf{c} = f(\mathbf{y}, \boldsymbol{\theta})$



input image and segmentation



3D representation of RGB values

Coding: Download Data and Setup MATLAB

The following data sets will be used throughout the course (and homework projects).

The following are ordered from small and easy to large and challenging:

- ▶ MNIST
- ▶ CIFAR-10
- ▶ CamVid: download from Mathworks web page.

References

- [1] Y. Bengio et al. Learning deep architectures for AI. *Foundations and trends® in Machine Learning*, 2(1):1–127, 2009.
- [2] A. Bordes, S. Chopra, and J. Weston. Question Answering with Subgraph Embeddings. *arXiv preprint arXiv:1406.3676*, 2014.
- [3] R. Collobert, J. Weston, L. Bottou, M. Karlen, K. Kavukcuoglu, and P. Kuksa. Natural Language Processing (Almost) from Scratch. *Journal of Machine Learning Research*, 12:2493–2537, 2011.
- [4] G. Hinton, L. Deng, D. Yu, G. E. Dahl, A.-r. Mohamed, N. Jaitly, A. Senior, V. Vanhoucke, P. Nguyen, T. N. Sainath, et al. Deep neural networks for acoustic modeling in speech recognition: The shared views of four research groups. *IEEE Signal Processing Magazine*, 29(6):82–97, 2012.
- [5] S. Jean, K. Cho, R. Memisevic, and Y. Bengio. On Using Very Large Target Vocabulary for Neural Machine Translation. *arXiv preprint arXiv:1412.2007*, 2014.
- [6] A. Krizhevsky, I. Sutskever, and G. Hinton. Imagenet classification with deep convolutional neural networks. *Advances in neural information processing systems*, 61:10971105, 2012.
- [7] Y. LeCun, Y. Bengio, and G. Hinton. Deep learning. *Nature*, 521(7553):436–444, 2015.

References (cont.)

- [8] Y. LeCun, B. E. Boser, and J. S. Denker. Handwritten digit recognition with a back-propagation network. In *Advances in neural information processing systems*, pages 396–404, 1990.
- [9] R. Raina, A. Madhavan, and A. Y. Ng. Large-scale deep unsupervised learning using graphics processors. In *the 26th Annual International Conference*, pages 873–880. ACM, June 2009.
- [10] F. Rosenblatt. The perceptron: A probabilistic model for information storage and organization in the brain. *Psychological review*, 65(6):386–408, 1958.
- [11] D. Rumelhart, G. Hinton, and J. Williams, R. Learning representations by back-propagating errors. *Nature*, 323(6088):533–538, 1986.