


# Single-Layer Neural Networks


Numerical Methods for Deep Learning

# Motivation: Nonlinear Models

In general, impossible to find a linear separator between classes



input features




transformed features

## Goal/Trick


Embed the points in higher dimension and/or move the points to make them linearly separable

# Motivation: Nonlinear Models

In general, impossible to find a linear separator between classes



input features



transformed features

## Goal/Trick

Embed the points in higher dimension and/or move the points to make them linearly separable

## Example: Linear Fitting

Assume  $\mathbf{C} \in \mathbb{R}^{n_c \times n}$ ,  $\mathbf{Y} \in \mathbb{R}^{n_f \times n}$  and  $n \gg n_f$ . Goal: Find  $\mathbf{W} \in \mathbb{R}^{n_c \times n_f}$  such that


$$\mathbf{C} = \mathbf{W}\mathbf{Y}$$

If  $\text{rank}(\mathbf{Y}) < n$ , there may be no solution.

Two options:

1. Regression: Solve  $\min_{\mathbf{W}} \|\mathbf{W}\mathbf{Y} - \mathbf{C}\|_F^2 \leadsto$  always has solutions, but residual might be large
2. Nonlinear Model: Replace  $\mathbf{Y}$  by  $\sigma(\mathbf{K}\mathbf{Y})$  in regression, where  $\sigma$  is element-wise function (aka activation) and  $\mathbf{K} \in \mathbb{R}^{m \times n_f}$  where  $m \gg n_f$

# Illustrating Nonlinear Models



## Remarks

- ▶ instead of  $\mathbf{W}\mathbf{Y} = \mathbf{C}$  solve  $\hat{\mathbf{W}}\sigma(\mathbf{K}\mathbf{Y}) = \mathbf{C}$
- ▶ solve bigger problem  $\leadsto$  memory, computation, ...
- ▶ what happens to  $\text{rank}(\sigma(\mathbf{K}\mathbf{Y}))$  when  $\sigma(x) = x$ ?

# Conjecture: Universal Approximation Properties

Given the data  $\mathbf{Y} \in \mathbb{R}^{n_f \times n}$  and  $\mathbf{C} \in \mathbb{R}^{n_c \times n}$  with  $n \gg n_f$ , there is nonlinear function  $\sigma : \mathbb{R} \rightarrow \mathbb{R}$ , a matrix  $\mathbf{K} \in \mathbb{R}^{m \times n_f}$ , and a bias  $\mathbf{b} \in \mathbb{R}^m$  such that

$$\text{rank}(\sigma(\mathbf{KY} + \mathbf{b})) = n.$$

Therefore, possible ?? to find  $\mathbf{W} \in \mathbb{R}^{n_c \times m}$

$$\mathbf{W}\sigma(\mathbf{KY} + \mathbf{b}) = \mathbf{C}.$$

# Choosing Nonlinear Model

$$\mathbf{W}\sigma(\mathbf{K}\mathbf{Y} + \mathbf{b}) = \mathbf{C}$$

- ▶ how to choose  $\sigma$ ?
  - ▶ early days: motivated by neurons
  - ▶ popular choice:  $\sigma(x) = \tanh(x)$  (smooth, bounded, ...)
  - ▶ nowadays:  $\sigma(x) = \max(x, 0)$  (aka ReLU, rectified linear unit, non-differentiable, not bounded, simple)
- ▶ how to choose  $\mathbf{K}$  and  $\mathbf{b}$ ?
  - ▶ pick randomly  $\leadsto$  branded as *extreme learning machines* ?
  - ▶ train (optimize)  $\leadsto$  done for most neural network
  - ▶ *deep learning* when neural network has many layers

# First Experiment: Random Transformation

Select activation function and choose  $\mathbf{K}$  and  $\mathbf{b}$  randomly and solve the least-squares/classification problem

The Pros:

- ▶ universal approximation theorem: can interpolate any function
- ▶ very(!) easy to program
- ▶ can serve as a benchmark to more sophisticated methods

Some concerns:

- ▶ may require very large  $\mathbf{K}$  (scale with  $n$ , number of examples)
- ▶ may not generalize well
- ▶ large dense linear algebra

EELM\_Peaks.m



# Learning the Weights

Assume that the number of examples,  $n$ , is very large.  
Using random weights,  $\mathbf{K}$  might need to be very large to fit training data.

Solution may not generalize well to test data.

Idea: Learn  $\mathbf{K}$  and  $b$  from the data (in addition to  $\mathbf{W}$ )


$$\min_{\mathbf{K}, \mathbf{W}, b} E(\mathbf{W}\sigma(\mathbf{K}\mathbf{Y} + \mathbf{b}), \mathbf{C}^{\text{obs}}) + \lambda R(\mathbf{W}, \mathbf{K}, \mathbf{b})$$

About this optimization problem:

- ▶ more unknowns  $\mathbf{K} \in \mathbb{R}^{m \times n_f}$ ,  $\mathbf{W} \in \mathbb{R}^{n_c \times m}$ ,  $\mathbf{b} \in \mathbb{R}^m$
- ▶ non-convex problem  $\leadsto$  local minima, careful initialization
- ▶ need to compute derivatives w.r.t.  $\mathbf{K}, \mathbf{b}$

# Non-Convexity

The optimization problem is non-convex. Simple illustration of cross-entropy along two random directions  $d\mathbf{K}$  and  $d\mathbf{W}$




(see `ESingleLayer_PlotObjective.m`)


Expect worse when number of layers grows!

# Training the Neural Network


- ▶ If non-convexity is not “too bad” can use standard gradient based methods
- ▶ If non-convexity is “ugly” need to modify standard methods (stochastic kick)
- ▶ If non-convexity is “bad” need global optimization techniques



good



bad



ugly

# Recap: Differentiating Linear Algebra Expressions

Easy ones:

$$F_1(\mathbf{x}, \mathbf{y}) = \mathbf{x}^\top \mathbf{y}$$

$$\mathbf{J}_x F_1(\mathbf{x}, \mathbf{y}) = \mathbf{y}^\top$$

$$F_2(\mathbf{A}, \mathbf{x}) = \mathbf{A}\mathbf{x}$$

$$\mathbf{J}_x F_2(\mathbf{x}, \mathbf{y}) = \mathbf{A}$$

How about

$$F_3(\mathbf{A}, \mathbf{X}) = \mathbf{A}\mathbf{X} \quad \mathbf{J}_{\text{vec}(\mathbf{X})} F_3 = ???$$

Recall that

$$\text{vec}(\mathbf{A}\mathbf{X}) = \text{vec}(\mathbf{A}\mathbf{X}\mathbf{I}) = (\mathbf{I} \otimes \mathbf{A})\text{vec}(\mathbf{X})$$

Therefore:

$$\mathbf{J}_{\text{vec}(\mathbf{X})} F_3(\mathbf{A}, \mathbf{X}) = \mathbf{I} \otimes \mathbf{A}$$

**Efficient mat-vec:**  $\mathbf{J}_{\text{vec}(\mathbf{X})} F \mathbf{v} = \text{vec}(\mathbf{A} \text{ mat}(\mathbf{v}))$

# Training Single Layer Neural Network

Assume no regularization (easy to add) and re-write optimization problem as

$$\min_{\mathbf{W}, \mathbf{K}, b} E(\mathbf{C}^{\text{obs}}, \mathbf{Z}, \mathbf{W}) \quad \text{with} \quad \mathbf{Z} = \sigma(\mathbf{K}\mathbf{Y} + b)$$

Agenda:

1. compute derivative of  $\text{vec}(\mathbf{Z})$  w.r.t.  $\text{vec}(\mathbf{K})$ ,  $b$
2. use chain rule to get

$$\mathbf{J}_{\text{vec}(\mathbf{K})} E = \mathbf{J}_{\text{vec}(\mathbf{Z})} E(\mathbf{C}^{\text{obs}}, \mathbf{Z}, \mathbf{W}) \mathbf{J}_{\text{vec}(\mathbf{K})} \mathbf{Z}$$

$$\mathbf{J}_b E = \mathbf{J}_{\text{vec}(\mathbf{Z})} E(\mathbf{C}^{\text{obs}}, \mathbf{Z}, \mathbf{W}) \mathbf{J}_b \mathbf{Z}$$

3. efficient code for mat-vecs with  $\mathbf{J}$  and  $\mathbf{J}^\top$

# Computing Jacobians

$$\mathbf{Z} = \sigma(\mathbf{KY} + b)$$

Recall that  $\sigma$  is applied element-wise.

$$\mathbf{J}_{\text{vec}(\mathbf{K})}\mathbf{Z} = \text{diag}(\sigma'(\mathbf{KY} + b))(\mathbf{Y}^\top \otimes \mathbf{I})$$

Efficient way to get matrix vector products

$$\begin{aligned}\mathbf{J}_{\text{vec}(\mathbf{K})}\mathbf{Z}\mathbf{v} &= \text{diag}(\sigma'(\mathbf{KY} + b))(\mathbf{Y}^\top \otimes \mathbf{I})\mathbf{v} \\ &= \text{vec}(\sigma'(\mathbf{KY} + b) \odot (\text{mat}(\mathbf{v})\mathbf{Y}))\end{aligned}$$

And for transpose get

$$\begin{aligned}(\mathbf{J}_{\text{vec}(\mathbf{K})}\mathbf{Z})^\top \mathbf{u} &= (\mathbf{Y} \otimes \mathbf{I})\text{diag}(\sigma'(\mathbf{KY} + b))\mathbf{u} \\ &= \text{vec}(\sigma'(\mathbf{KY} + b) \odot \text{mat}(\mathbf{u})\mathbf{Y}^\top)\end{aligned}$$

# Class Problems: Derivatives of Single Layer

## Derivations:

1. Compute  $\mathbf{J}_b \mathbf{Z} \mathbf{v}$  and  $(\mathbf{J}_b \mathbf{Z})^\top \mathbf{u}$
2. Compute  $\mathbf{J}_{\text{vec}(\mathbf{Y})} \mathbf{Z} \mathbf{v}$  and  $(\mathbf{J}_{\text{vec}(\mathbf{Y})} \mathbf{Z})^\top \mathbf{u}$

## Coding:

```
function[Z,JKt,Jbt,JYt,JK,Jb,JY] = singleLayer(K,b,Y)
% Returns Z = sigma(K*Y+b) and
%                               functions for J'*U and J*V
```

## Testing:

1. Derivative check for Jacobian mat-vec
2. Adjoint tests for transpose, let  $\mathbf{v}, \mathbf{u}$  be arbitray vectors

$$\mathbf{u}^\top \mathbf{J} \mathbf{v} \approx \mathbf{v}^\top \mathbf{J}^\top \mathbf{u}$$

# Putting Things Together

Implement loss function of single-layer NN

$$E(\mathbf{K}, b, \mathbf{W}) \stackrel{\text{def}}{=} E(\mathbf{C}, \mathbf{Z}, \mathbf{W}), \quad \mathbf{Z} = \sigma(\mathbf{K}\mathbf{Y} + b)$$

```
function [Ec,dE] = singleLayerNNObjFun(x,Y,C,m)
% where x = [K(:); b; W(:)]
% evaluates single layer and computes cross entropy
%           and gradient (extend for approx. Hessian)
```

Use

1.  $\nabla_{\mathbf{Z}} E = \mathbf{W}^{\top} \nabla_{\mathbf{S}} E(\mathbf{S}), \quad \mathbf{S} = \mathbf{WZ}$
2.  $\nabla_{\mathbf{K}} E = \mathbf{J}_{\mathbf{K}}^{\top} \nabla_{\mathbf{Z}} E$
3.  $\nabla_b E = \mathbf{J}_b^{\top} \nabla_{\mathbf{Z}} E$
4.  $\nabla_{\mathbf{W}} E = \nabla_{\mathbf{S}} E(\mathbf{S}) \mathbf{Y}$



# Test Problem

Before going to real data, let us try the *inverse crime*.  
Generate data

```
n    = 500; nf = 50; nc = 10; m    = 40;  
Wtrue = randn(nc,m);  
Ktrue = randn(m,nf);  
btrue = .1;  
  
Y      = randn(nf,n);  
Cobs   = exp(Wtrue*singleLayer(Ktrue,btrue,Y));  
Cobs   = Cobs./sum(Cobs,1);
```

Goal: Reconstruct Wtrue, Ktrue, btrue!

# Gauss-Newton Method

**Goal:** Use curvature information for fast convergence

$$\nabla_{\mathbf{K}} E(\mathbf{K}, \mathbf{b}, \mathbf{W}) = (\mathbf{J}_{\mathbf{K}} \mathbf{Z})^{\top} \nabla_{\mathbf{Z}} E(\mathbf{W} \sigma(\mathbf{K} \mathbf{Y} + \mathbf{b}), \mathbf{C}),$$

where  $\mathbf{J}_{\mathbf{K}} \mathbf{Z} = \nabla_{\mathbf{K}} \sigma(\mathbf{K} \mathbf{Y} + \mathbf{b})^{\top}$ . This means that Hessian is

$$\begin{aligned} \nabla_{\mathbf{K}}^2 E(\mathbf{K}) &= (\mathbf{J}_{\mathbf{K}} \mathbf{Z})^{\top} \nabla_{\mathbf{Z}}^2 E(\mathbf{C}, \mathbf{Z}, \mathbf{W}) \mathbf{J}_{\mathbf{K}} \mathbf{Z} \\ &\quad + \sum_{i=1}^n \sum_{j=1}^m \nabla_{\mathbf{K}}^2 \sigma(\mathbf{K} \mathbf{Y} + \mathbf{b})_{ij} \nabla_{\mathbf{Z}} E(\mathbf{C}, \mathbf{Z}, \mathbf{W})_{ij} \end{aligned}$$

First term is spsd and we can compute it.

We neglect second term since

- ▶ can be indefinite and difficult to compute
- ▶ small if transformation is roughly linear or close to solution (easy to see for least-squares)

do the same for  $\mathbf{b}$  and use full Hessian for  $\mathbf{W} \rightsquigarrow$  ignore coupling!

## Experiment: Adversarial Example

Suppose you have trained your network  $\rightsquigarrow \mathbf{K}, b, \mathbf{W}$  so that validation loss is low. This means that for most examples  $\mathbf{y}$ ,

$$\mathbf{W}\sigma(\mathbf{K}\mathbf{y} + b) \approx \mathbf{c}.$$

An adversary might try to fool this classifier by adding a small perturbation  $\mathbf{d}$  to the example to achieve a desired label  $\hat{\mathbf{c}}$ .

Formulate as optimization problem

$$\min_{\mathbf{d}} E(\mathbf{W}\sigma(\mathbf{K}(\mathbf{y} + \mathbf{d}) + b), \hat{\mathbf{c}})$$

- ▶ setup objective function
- ▶ think about constraints, regularization